


F2 Higher concurrency can lower throughput

70B · color = quantization · line / marker = input length



Two RTX PRO 6000 Blackwell GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens
One run per measured setting; missing and capacity-excluded values are unplotted. Small gaps do not establish a winner.