


F3 Throughput versus concurrency

70B · bar labels: tok/s · top labels: change from C8 to C16



Two RTX PRO 6000 Blackwell GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens
One run per measured setting; missing and capacity-excluded values are unplotted. Small gaps do not establish a winner.