


F4 Context changes quantization gains

70B · concurrency = 8 · IQ1_M versus Q4_K_M



Two RTX PRO 6000 Blackwell GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

One run per measured setting; missing and capacity-excluded values are unplotsed. Small gaps do not establish a winner.