


Runtime Cost · Llama 70B

Four views of measured 70B serving · IQ1_M, Q2_K, Q4_K_M and Q8_0

Available measurements: 36/60 settings · Missing and capacity-excluded values are unplotted; see coverage.csv.

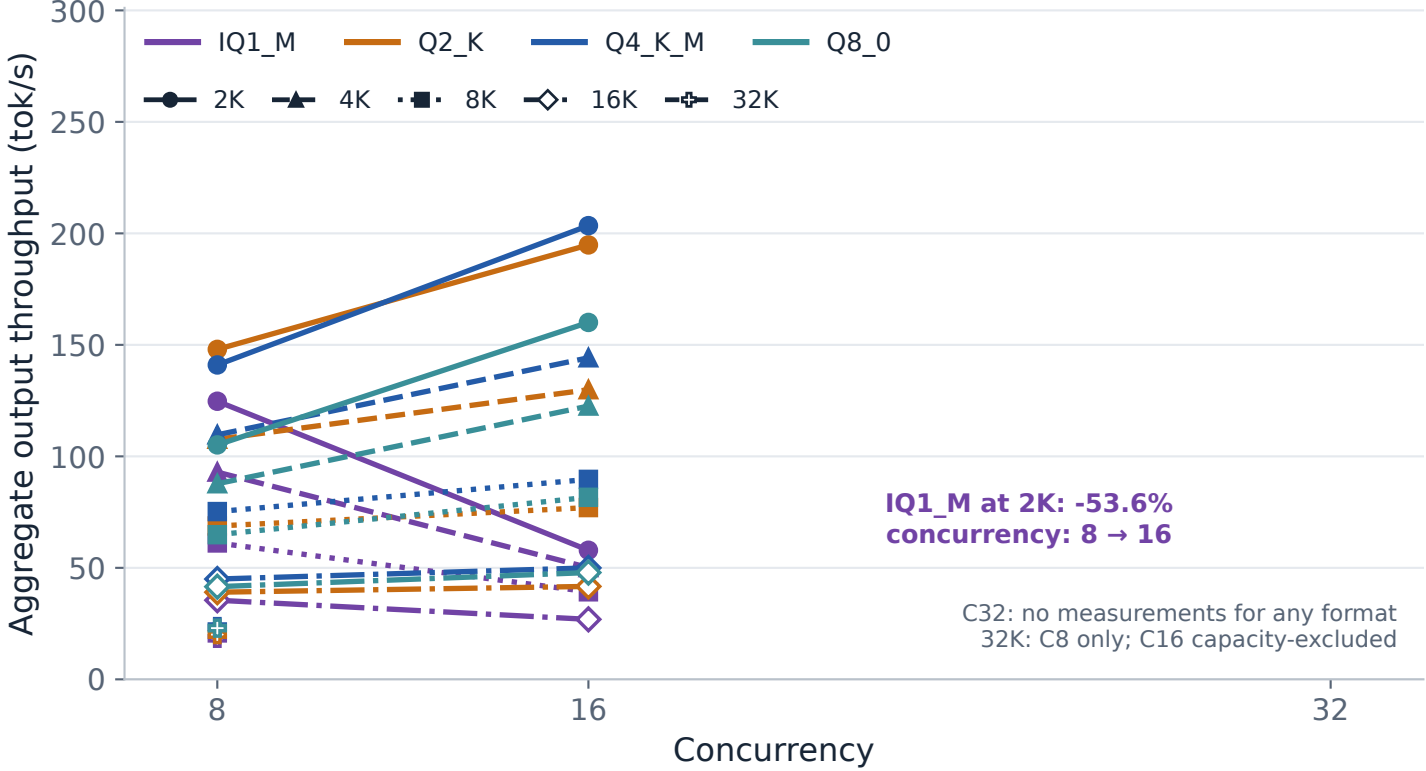
F1 Lower bits ≠ lower latency

70B · concurrency = 8 · percentages relative to Q4_K_M



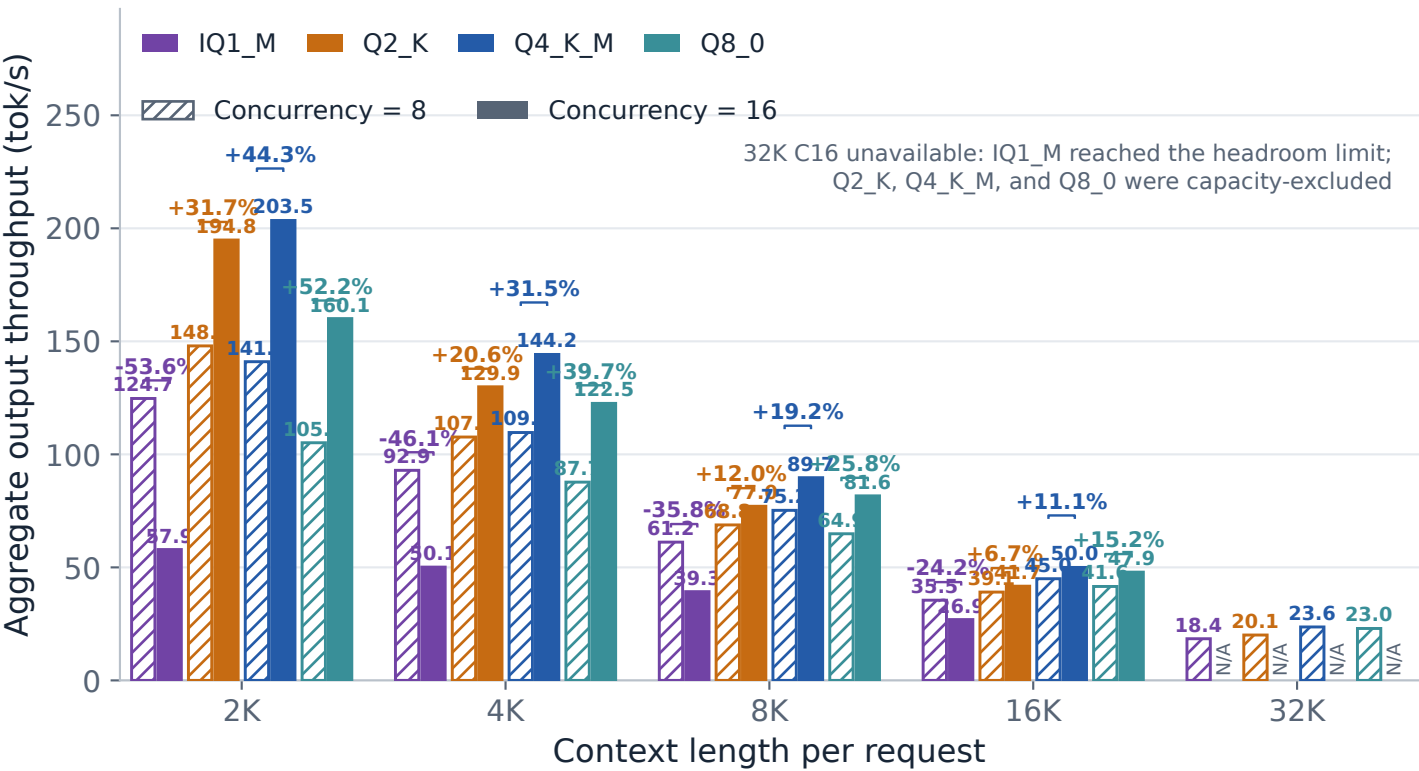
F2 Higher concurrency can lower throughput

70B · color = quantization · line / marker = input length



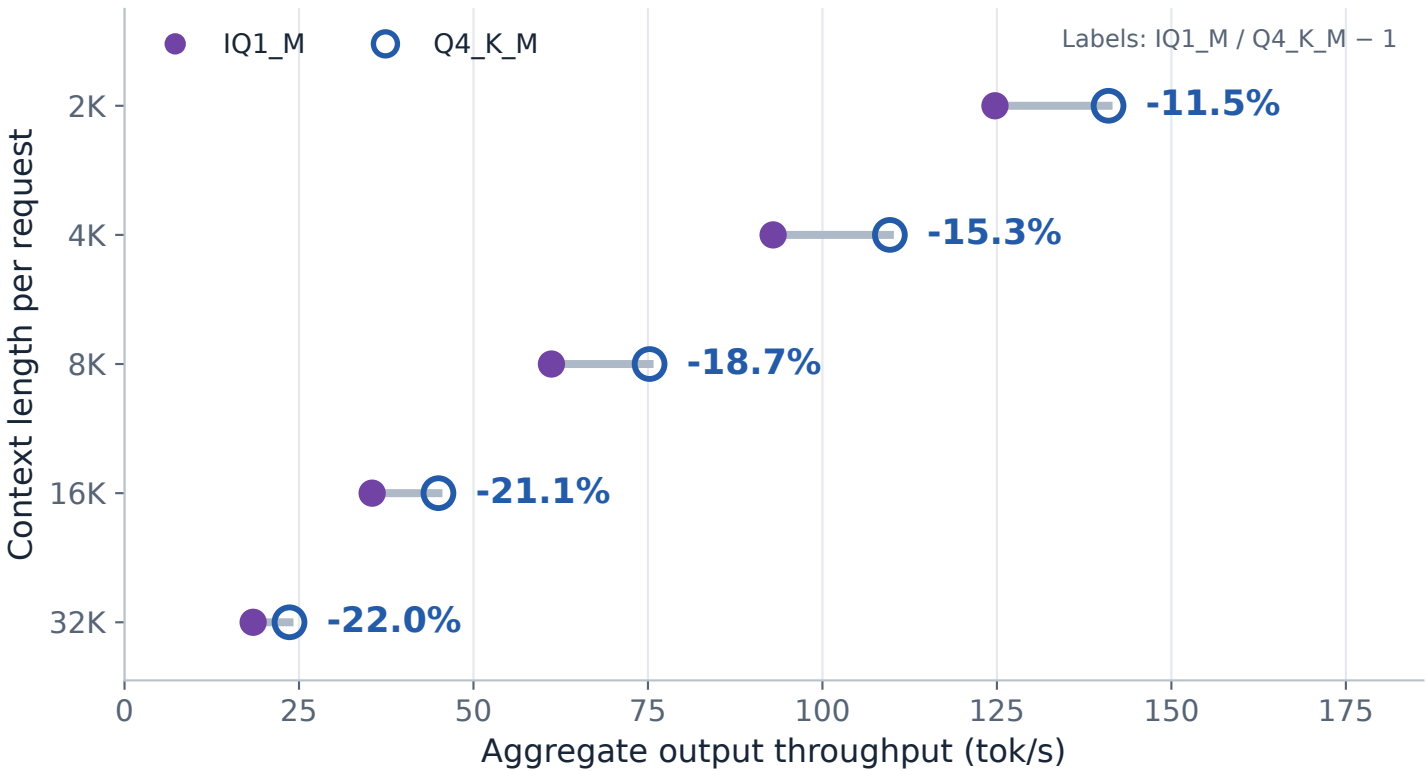
F3 Throughput versus concurrency

70B · bar labels: tok/s · top labels: change from C8 to C16



F4 Context changes quantization gains

70B · concurrency = 8 · IQ1_M versus Q4_K_M



Two RTX PRO 6000 Blackwell GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens
One run per measured setting; missing and capacity-excluded values are unplotted. Small gaps do not establish a winner.