

# Available context measurements

70B · largest validated input within the 2K–16K grid

| Quantization | Concurrency           |                     |                     |
|--------------|-----------------------|---------------------|---------------------|
|              | 8                     | 16                  | 32                  |
|              | IQ1_M -<br><b>16K</b> | <b>8K</b>           | <b>Not measured</b> |
|              | Q2_K -<br><b>16K</b>  | <b>8K</b>           | <b>4K</b>           |
| Q4_K_M -     | <b>16K</b>            | <b>8K</b>           | <b>4K</b>           |
| Q8_0 -       | <b>Not measured</b>   | <b>Not measured</b> | <b>Not measured</b> |

Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

One run per setting; thermal limiting occurred. Small gaps do not establish a winner.