


F1 Lower bits $\neq$ lower latency

70B · concurrency = 8 · percentages relative to Q4_K_M



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

One run per setting; thermal limiting occurred. Small gaps do not establish a winner.