


F1 Lower bits $\neq$ lower latency

70B · 8 clients · paired Q2_K / Q4_K_M comparison



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

Measured points: one run per setting; thermal limiting occurred in some settings.