

F2 More concurrency $\neq$ more throughput

8B · IQ1_M · input lengths shown as separate lines



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

Measured points: one run per setting; thermal limiting occurred in some settings.