


F2 Higher concurrency can lower throughput

70B · color = quantization · line / marker = input length



**IQ1_M at 2K: -39.1%
concurrency: 8 → 16**

IQ1_M: C32 not measured
Q8_0: no measurements
16K: C8 only; other gaps listed in coverage.csv

Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens
One run per setting; thermal limiting occurred. Small gaps do not establish a winner.