


F2-F4 Throughput across prompt lengths

Llama 70B only · color = quantization · line / marker = concurrency



F2 IQ1_M at 2K, concurrency = 8 / 16 / 32: 49.4 tok/s / 30.1 tok/s / not measured

F3 Q4_K_M gain from concurrency 8 → 16: +81.9% at 2K, +52.9% at 4K, +30.2% at 8K

F4 IQ1_M / Q4_K_M throughput change at concurrency = 8: +7.6% at 2K, -0.2% at 4K, -6.7% at 8K, -12.1% at 16K

Q8_0: no measured values at concurrency = 8, 16 or 32. Missing/capacity-excluded values are unplotted.

Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

One run per setting; thermal limiting occurred. Small gaps do not establish a winner.