


F3 Long context weakens concurrency scaling

70B · Q4_K_M · throughput gain from 8 → 16 clients



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

Measured points: one run per setting; thermal limiting occurred in some settings.