


# F4 Context changes quantization gains

70B · concurrency = 8 · IQ1\_M versus Q4\_K\_M



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

One run per setting; thermal limiting occurred. Small gaps do not establish a winner.