


F4 Long context erodes quantization speedup

8B · 32 clients · FP16 versus Q4_K_M



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens

Measured points: one run per setting; thermal limiting occurred in some settings.