


Longer context reduces quantization gains

Llama 8B · Q4_K_M throughput advantage over FP16



Gain = $100 \times (\text{Q4_K_M throughput} / \text{FP16 throughput} - 1)$; matched context and concurrency.

Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens
One run per setting; thermal limiting occurred. Small gaps do not establish a winner.