


F1 Lower bits ≠ lower latency

8B · concurrency = 8 · percentages relative to Q4_K_M



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens
One run per setting; thermal limiting occurred. Small gaps do not establish a winner.