


F3 Throughput versus concurrency


Llama 8B · point labels: tok/s · percentages: change relative to concurrency = 8

● 2K context ▲ 4K context ■ 8K context ◆ 16K context


IQ1_M



Q2_K



Q4_K_M



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens
One run per setting; thermal limiting occurred. Small gaps do not establish a winner.