


Runtime Cost

Four findings from quantized Llama serving

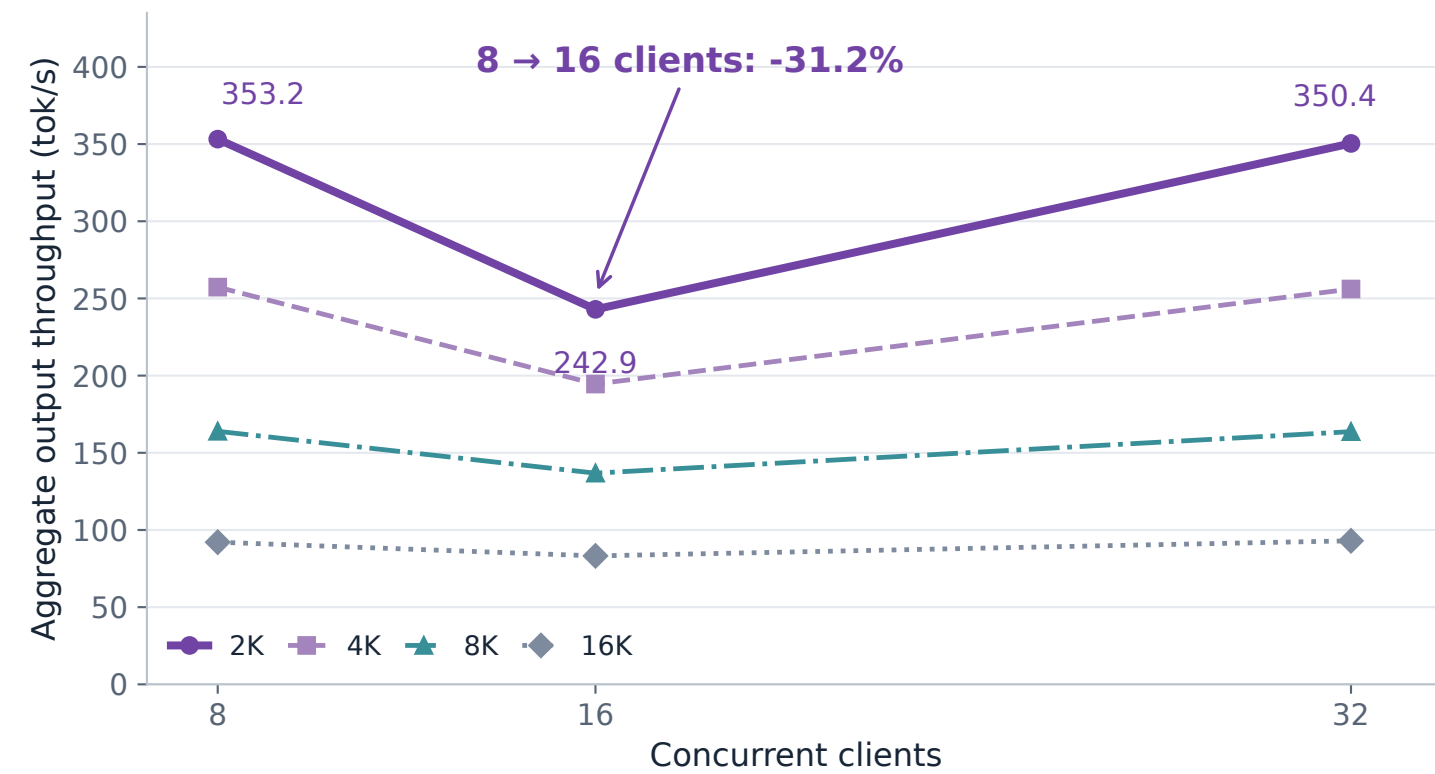
F1 Lower bits $\neq$ lower latency

70B · 8 clients · paired Q2_K / Q4_K_M comparison



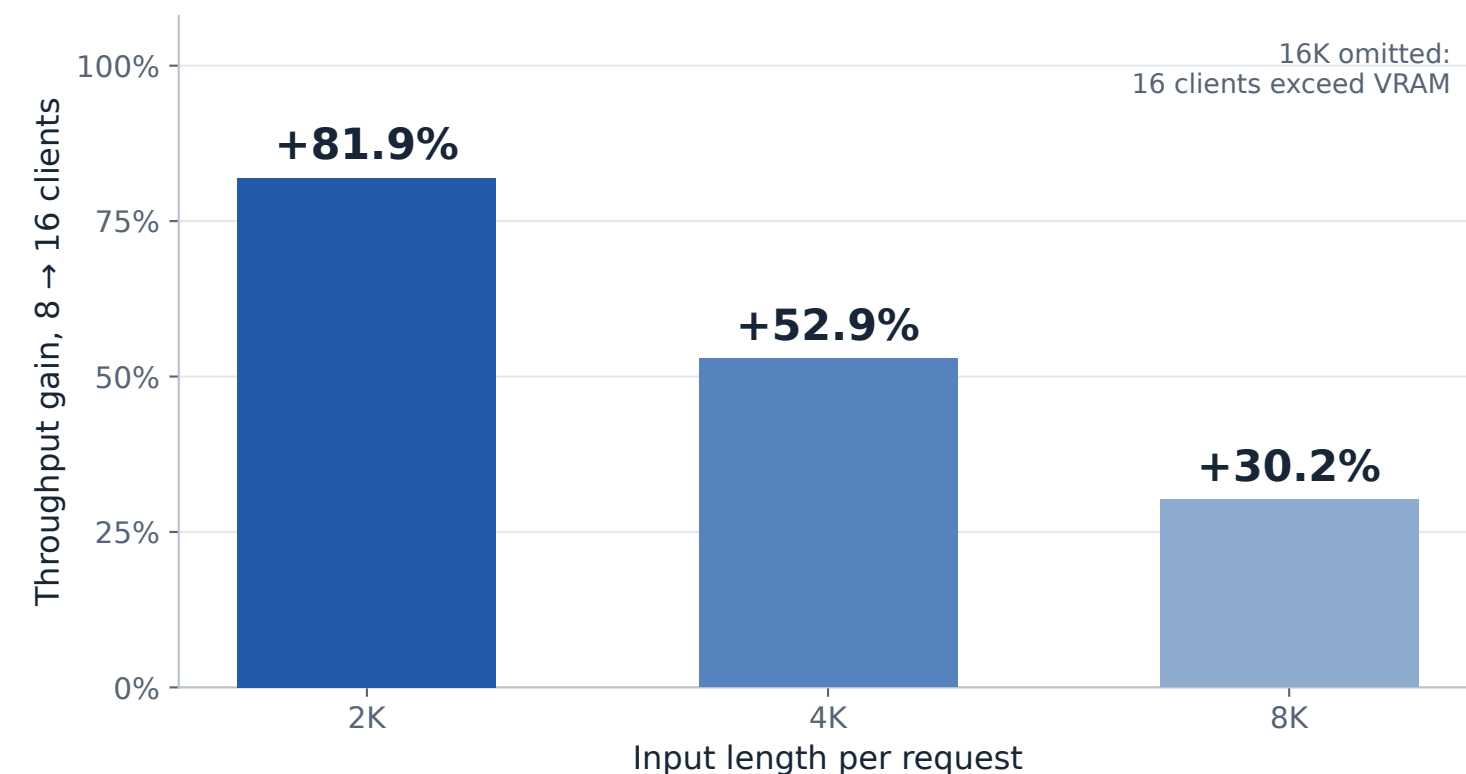
F2 More concurrency $\neq$ more throughput

8B · IQ1_M · input lengths shown as separate lines



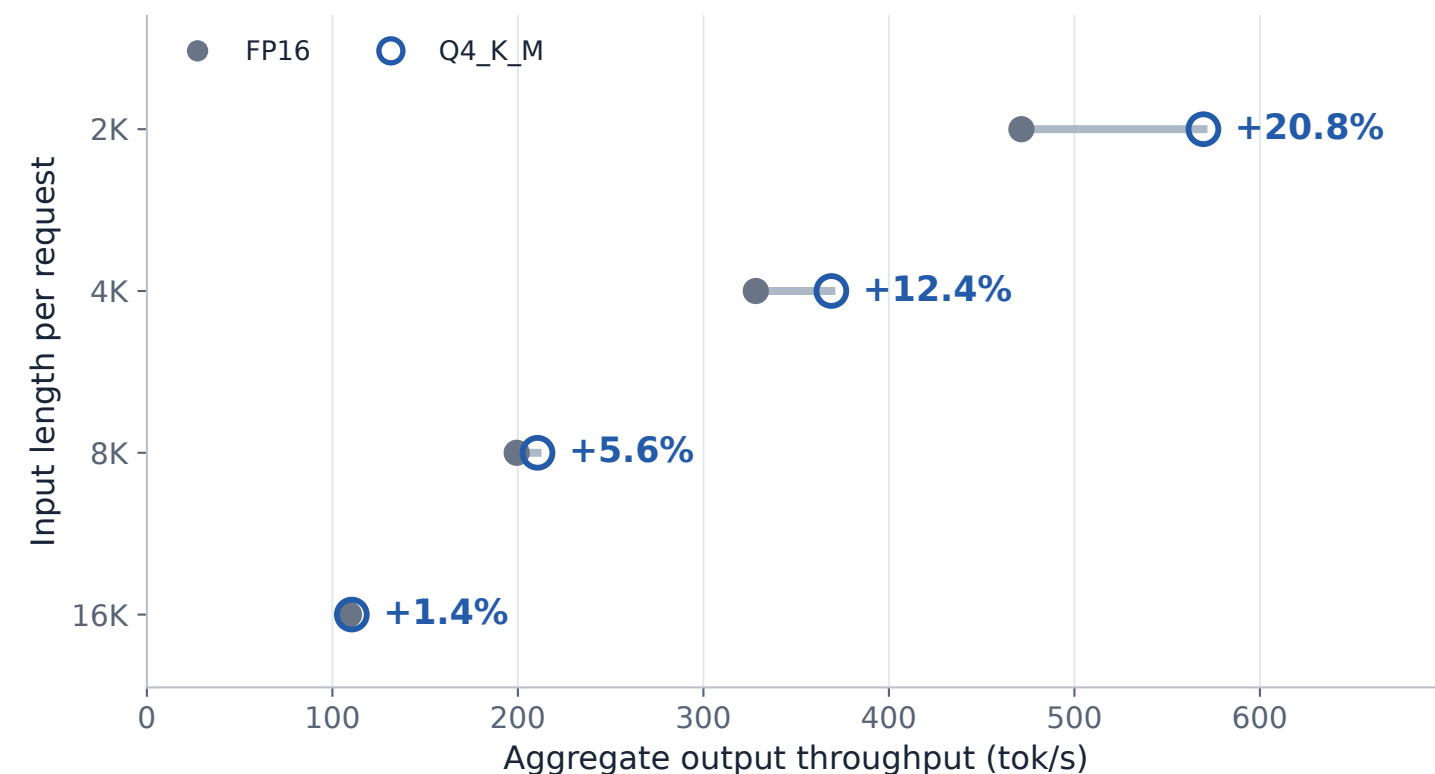
F3 Long context weakens concurrency scaling

70B · Q4_K_M · throughput gain from 8 → 16 clients



F4 Long context erodes quantization speedup

8B · 32 clients · FP16 versus Q4_K_M



Two RTX A6000 GPUs · FP16 KV · 512 outputs/request · prompt reuse off · K = 1,024 input tokens
Measured points: one run per setting; thermal limiting occurred in some settings.