







Decode: more instruction work and register pressure

70B gate/up · GPU 0 · N = 8 · median of five matched launches

A Registers, block limits, and achieved occupancy



B Instruction work, memory traffic, and kernel duration



Actual values on separate, zero-based scales. Q2 labels show percentage change relative to Q4.

Higher register use and lower occupancy accompany the slowdown; the independent latency cost of occupancy is not isolated.