


Other kernels · Q4 and Q2 on a common roofline


Scalar FP32 view · 70B · 8 clients · 2K input / 512 output

Color / shape identifies the operator; filled markers are Q4 and hollow markers are Q2.

Prefill · Q4_K_M and Q2_K



Decode · Q4_K_M and Q2_K



- Q4_K_M: filled
- Q2_K: hollow
- DRAM ceiling
- FP32 ceiling
- ... Compute / memory boundary

GPU0 points; † Q4 LM head is GPU1. FlashAttention retains all three decode launch configurations per format.

31 operator/configuration groups from 120 launches. Coincident points retain their measured coordinates; labels use leader lines.

Scalar FP32 counts exclude Tensor Core, FP16, integer and special-function work. Missing kernels are listed in the report.