


Other kernels · scalar FP32 roofline

70B · 8 clients · 2K input / 512 output · filled Q4 / hollow Q2

Prefill · Q4_K_M and Q2_K



GPU0; † Q4 decode LM head is GPU1. Each point retains one operator/configuration; no missing coordinates are fabricated.

Scalar FP32 counts exclude Tensor Core, FP16, integer and special-function work. Rated per-GPU ceilings: 38.7 TFLOP/s, 768 GB/s.